

Multimodal VQA Architecture

CLIP (Frozen) + Trainable Projection + Question-Type-Specific Heads

